

How I Learned to Stop Worrying and Love the R Console

Marathon Edition

Irfan E Kanat

Department of Digitization
College of Business
Ohio University

February 20, 2019



Outline

- 1 Introduction
- 2 R Console
- 3 Importing Data
- 4 Packages
- 5 Sample Analysis and Visualizations
 - Descriptive Visualizations
 - Modeling
- 6 Reporting
- 7 Where to Go Next?

Who am I?

Irfan Kanat

R user since 2006

Open Source Evangelist

Before We Begin

Got R & R Studio Installed?

Get your workshop documents:

<https://github.com/iekanat/RWorkshopMarathon>

What is this about?

A brief introduction to R.

- R Console
- Importing Data
- Packages
- Sample analyses
- Where to get help?



What is R?

From R project web site:

R is a language *and* an environment for statistical computing and graphics.

- Language
- Environment
- Statistics and Visualization



What is R?

All this means R is very flexible, which played a huge role in its success.

My take: Low cost, high quality, open source solution for your analysis needs.

When to Use R?


R is very strong for your classical machine learning and statistical analysis. Thousands of packages address almost all analysis needs. It is a logical first step to start analysis.

When to Use R?

R is very strong for your classical machine learning and statistical analysis. Thousands of packages address almost all analysis needs. It is a logical first step to start analysis.

Yet it's core design is starting to show its age. There are certain down sides to traditional R:

- Everything is stored in memory¹
- R is single core¹

¹Except when it is not. There are packages to overcome these issues. 

Best Part of R

Packages CRAN houses over 13K packages. Providing functionality way beyond what is available in commercial packages.

Community Millions of users mean, all your questions are either already answered or will be in hours.

Performance While memory and core restrictions are real, for the cost of a single user license of a commercial package, you can buy better hardware to run R. Furthermore, with the packages providing multicore and flatfile functionality, R performance is on par or better than commercial packages

Outline

- 1 Introduction
- 2 R Console
- 3 Importing Data
- 4 Packages
- 5 Sample Analysis and Visualizations
 - Descriptive Visualizations
 - Modeling
- 6 Reporting
- 7 Where to Go Next?

Command Driven Interface

Command line may be intimidating

Power over Convenience

Consider the number of

- Functions
- Parameters
- Data sources
- Variables
- Replications


Command Driven Interface

Command line may be intimidating

Power over Convenience

Consider the number of

- Functions
- Parameters
- Data sources
- Variables
- Replications



R Studio

The screenshot displays the R Studio interface with four main panels:

- Source Editor (Top Left):** Contains R code for a project setup and a simple plot. A red box highlights the code, and a large red letter 'A' is overlaid on it.


```

30 # Create a new R project directory (i.e. a folder that you place your project directory
31 # to a place you can easily access and remember (like Documents folder or Home folder).
32
33 # Create a new R project directory
34
35 ## Getting Your Feet Wet
36
37 R console is quite flexible. You can use it for a number of purposes. Below are some basic algebra
38 examples.
39
40 # Create a new R project directory
41 # Create a new R project directory
42 # Create a new R project directory
43 # Create a new R project directory
44 # Create a new R project directory
      
```
- Environment (Top Right):** Shows the Global Environment with variables:
 - Data:**
 - admtData: 400 obs. of 4 variables
 - csv: 24 obs. of 3 variables
 - d: 10 obs. of 3 variables
 - xls: 29 obs. of 3 variables
 - Values:**
 - a: 5
 - b: int [1:30] 1 2 3 4 5 6 7 8 9 10 ...
 - c: num [1:30] 6 7 8 9 10 11 12 13 14 15 ...
 - country: Large SpatialPolygonsDataFrame (244 elements, 1.5 Mb)
 - fac: chr [1:10] "C" "B" "C" "C" "B" "A" "C" "A" "C" "A"
 - logit_0: List of 30
 - svm_0: List of 30
 - Functions:**
 - hello: function (x)
 A large blue letter 'C' is overlaid on the Environment panel.
- Console (Bottom Left):** Shows the output of the R code, including model deviance and Fisher Scoring iterations. A large green letter 'B' is overlaid on the console output.


```

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 499.98 on 399 degrees of freedom
Residual deviance: 458.52 on 394 degrees of freedom
AIC: 470.52

Number of Fisher Scoring iterations: 4

> plot(svm_0, admtData, gre-gpa) # Let us plot the results
Error in xy.coords(x, y, xlabel, ylabel, log):
  'x' and 'y' lengths differ
> svm_0<-svm(admt--, data=admtData, type="C-classification")
Error: could not find function "svm"
> library(e1071) # Load package into work space.
> svm_0<-svm(admt--, data=admtData, type="C-classification")
> plot(svm_0, admtData, gre-gpa) # Let us plot the results
> plot(svm_0, admtData, gre-gpa, split = list(rank="1")) # Let us plot the results
> read.c
      
```
- Plots (Bottom Right):** Displays an SVM classification plot titled "SVM classification plot". The x-axis is labeled "gpa" (ranging from 2.5 to 4.0) and the y-axis is labeled "gre" (ranging from 300 to 800). The plot shows two classes of data points (circles and crosses) separated by a decision boundary. A yellow circle highlights a specific data point. A color bar on the right indicates the predicted class (0 in cyan, 1 in magenta).

New Project

File > New Project

Empty Directory > Empty Project > Directory Name: Workshop

R as a Calculator I

```
# Arithmetics
```

```
2 + 2
```

```
## [1] 4
```

```
2 * 3
```

```
## [1] 6
```

```
2^3
```

```
## [1] 8
```

```
log(100, 10)
```

```
## [1] 2
```


R as a Calculator II

```
# Logic
1 == 2

## [1] FALSE

1 != 2

## [1] TRUE

2 < 3

## [1] TRUE
```

Variables I

```
A <- 2
```

```
A
```

```
## [1] 2
```

```
# a # Case sensitive
```

```
"A" != "a" # Explanation
```

```
## [1] TRUE
```

```
B <- 7
```

```
A + B
```

```
## [1] 9
```

Variables II

```
C <- c(1, 3, 7, 9)  # A list can be in a variable
```

```
C
```

```
## [1] 1 3 7 9
```

```
C + A
```

```
## [1] 3 5 9 11
```

```
C * A
```

```
## [1] 2 6 14 18
```

```
C < 5
```

```
## [1] TRUE TRUE FALSE FALSE
```

Indexes and Data Frames I

```
1:30
```

```
## [1] 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23  
## [24] 24 25 26 27 28 29 30
```

```
C[3]
```

```
## [1] 7
```

```
C[c(2, 3)]
```

```
## [1] 3 7
```

```
C[1:3]
```

```
## [1] 1 3 7
```

Indexes and Data Frames II

```
Countries <- data.frame(names=c("US", "TR", "DE"), supply=c(10, 8, 7),  
                        those=c(TRUE, FALSE, FALSE))
```

Countries

```
##   names supply those  
## 1    US     10  TRUE  
## 2    TR      8 FALSE  
## 3    DE      7 FALSE
```

Indexes and Data Frames III

```
Countries[2, ]
```

```
##      names supply those  
## 2      TR          8 FALSE
```

```
Countries[, 3]
```

```
## [1] TRUE FALSE FALSE
```

```
Countries[2, 3]
```

```
## [1] FALSE
```

```
Countries[1:2, ]
```

```
##      names supply those  
## 1      US      10  TRUE  
## 2      TR       8  FALSE
```

Indexes and Data Frames IV

```
Countries[, "names"]
```

```
## [1] US TR DE  
## Levels: DE TR US
```

```
Countries$names
```

```
## [1] US TR DE  
## Levels: DE TR US
```

```
Countries$those
```

```
## [1] TRUE FALSE FALSE
```

Loops in R

CAUTION!

R is notoriously inefficient with your classic loops

- Structure of the Data Frame
- Memory Management

Try to use an apply function instead.

Vectorize your operations.

For Loop in R

```
for (i in 1:3) print(i)

## [1] 1
## [1] 2
## [1] 3

# Iterating through a data frame
for (i in 1:nrow(Countries)) {
  print(Countries[i, ])
}

##      names supply those
## 1      US      10  TRUE
##      names supply those
## 2      TR       8 FALSE
##      names supply those
## 3      DE       7 FALSE
```

Functions I

```
mean(C)  # Takes parameters
```

```
## [1] 5
```

```
mean(C, trim = 0.1, na.rm = T)  # Takes multiple parameters
```

```
## [1] 5
```

```
log(sum(C)/length(C))  # Can be combined
```

```
## [1] 1.609438
```

Functions II

```
HelloWorld <- function(x, y = 1) {  
  for (i in 1:y) {  
    print(paste("Hello", x))  
  }  
}
```

```
HelloWorld("COB Faculty")
```

```
## [1] "Hello COB Faculty"
```

```
HelloWorld("COB Faculty", 2)
```

```
## [1] "Hello COB Faculty"
```

```
## [1] "Hello COB Faculty"
```

Functions III

```
HelloWorld  # Review the source code
```

```
## function(x, y = 1) {
##   for (i in 1:y) {
##     print(paste("Hello", x))
##   }
## }
## <bytecode: 0x143953e8>
```

```
ls
```

```
## function (name, pos = -1L, envir = as.environment(pos), all.names = FALSE,
##   pattern, sorted = TRUE)
## {
##   if (!missing(name)) {
##     pos <- tryCatch(name, error = function(e) e)
##     if (inherits(pos, "error")) {
##       name <- substitute(name)
##       if (!is.character(name))
##         name <- deparse(name)
##       warning(gettextf("%s converted to character string",
##         sQuote(name))), domain = NA)
##     }
##   }
## }
```

Commonly Used Functions I

```
ls() # Get a list of objects in the workspace

## [1] "A"          "admitData"  "B"          "C"          "Countries"
## [6] "HelloWorld" "i"
```



```
ls(pattern = "*_0") # partial match on object search

## character(0)
```



```
rm("svm_0") # Remove an object from the workspace

## Warning in rm("svm_0"): object 'svm_0' not found

# rm(list=ls()) # This would remove everything if ran
```

Commonly Used Functions II

```
mean(A)    # Mean

## [1] 2

sd(admitData[, "gre"]) # Standard Deviation

## [1] 115.5165

# AIC()
```

Commonly Used Functions III

```
str(Countries) # Look at the structure of objects
```

```
## 'data.frame': 3 obs. of 3 variables:
## $ names : Factor w/ 3 levels "DE","TR","US": 3 2 1
## $ supply: num 10 8 7
## $ those : logi TRUE FALSE FALSE
```

```
summary(Countries) # Get summary of data
```

```
## names      supply      those
## DE:1   Min.    : 7.000   Mode :logical
## TR:1   1st Qu.: 7.500   FALSE:2
## US:1   Median : 8.000   TRUE :1
##          Mean    : 8.333
##          3rd Qu.: 9.000
##          Max.    :10.000
```

Commonly Used Functions IV

```
summary(logit_1) # Get summary of model


##
## Call:
## glm(formula = admit ~ gre, family = "binomial", data = admitData)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.1623  -0.9052  -0.7547   1.3486   1.9879
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -2.901344   0.606038  -4.787 1.69e-06 ***
## gre          0.003582   0.000986   3.633 0.00028 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 499.98  on 399  degrees of freedom
## Residual deviance: 486.06  on 398  degrees of freedom
## AIC: 490.06
##
## Number of Fisher Scoring iterations: 4
```


Commonly Used Functions V

```
cor(admitData[, 1:3]) # Get correlation matrix
```

```
##           admit           gre           gpa
## admit 1.0000000 0.1844343 0.1782123
## gre   0.1844343 1.0000000 0.3842659
## gpa   0.1782123 0.3842659 1.0000000
```

Questions



Outline

- 1 Introduction
- 2 R Console
- 3 Importing Data**
- 4 Packages
- 5 Sample Analysis and Visualizations
 - Descriptive Visualizations
 - Modeling
- 6 Reporting
- 7 Where to Go Next?

Importing Data

R allows importing data from a wide variety of sources.

- Comma Separated Values (CSV)
- Databases
- Flat files
- Lesser statistical packages
- and more

Importing CSV Files

CSV has certain advantages that make it popular.

- Compatibility
- Flexibility
- Simplicity

Sample

```
"iso2", "Supply", "Those"  
"AU", 20, 0  
"TR", 80, 1  
"US", 100, 0  
"GB", 50, 0  
"DE", 70, 0
```

We use `read.csv()` or `read.csv2()` commands to import the csv files.

```
saveData <- read.csv("PathToCSV", header = TRUE, sep = ",", quote = "\"")
```

Working with Excel Files

Much like CSV, except it lacks the simplicity, flexibility, and compatibility of CSV.

```
# Load the necessary library  
library(xlsx)  
# Read in the data from excel file  
xlsx <- read.xlsx("country.xlsx", sheetIndex = 1)
```

Working with Databases

No speed advantage.

Data larger than memory.

Working with databases:

- Work in the database.
- Import data from database.



Working with Databases

```
# Load the necessary library
library(RMySQL)
# Establish connection to the database.
channel <- dbConnect(MySQL(), user = "uname", password = "pwd", host = "127.0.0.1",
  dbname = "exampledata")
# Send query and save results in R workspace
sql <- dbGetQuery(channel, "SELECT * FROM table;")
```


Lesser Statistical Packages :P

Foreign Package

Newer file formats

- sas7bdat
- readstata13



Questions



Outline

- 1 Introduction
- 2 R Console
- 3 Importing Data
- 4 Packages**
- 5 Sample Analysis and Visualizations
 - Descriptive Visualizations
 - Modeling
- 6 Reporting
- 7 Where to Go Next?

Packages: Source of R's Power

Encountered already

Make R extendible

Like libraries

Collection of:

- functions
- documentation
- data files



Gifts from the Community

Currently over 13000 packages

for

- Statistical Modeling
- Machine Learning
- Data Manipulation
- Visualization
- ...

from

- Economics
- Computer Science
- Statistics
- Medicine
- ...



Great but Where are My Gifts?

Comprehensive R Archive Network (CRAN)

A Group of FTP and HTML servers hosting R packages.

R has built in package management facilities.

Most of these can be achieved through the R Studio GUI. (Area D, packages pane)

Package Management

```
# Installing a package
install.packages("e1071") # Notice the quotes around package name
# Loading package into memory
library(e1071) # Notice the lack of quotes
# Unload package
detach("package:e1071", unload = TRUE) # Notice the package: prepended

# Get the list of packages loaded
(.packages())
# Get list of all installed packages (output omitted)
.packages(all.available = T)
```

How to Find Packages

If you want to search a certain word in installed packages' documentation, you can always use `??` or `help.search()`

```
??mixed  
help.search("mixed model")
```

Internet searches are a bit problematic as R can be a bit ambiguous until Google learns you are interested in the statistical computing environment.

Comprehensive R Archive Network (CRAN)

R Forge

R site search also available with command `RSiteSearch()`

R seek

Commonly Used Packages: Data Manipulation

data.tables Replaces traditional `data.frame`.

- Faster access/write
- Improved selection
- Improved subsetting
- Improved aggregation

Not a drop-in replacement as it breaks compatibility in some cases.

ddplyr

Additional functionality for:

- selection
- filtering
- aggregation

Provides efficient back-end data structures to speed things up.

Works with databases as well.

Commonly Used Packages: Statistics

Multiple Regression: Stats package, `lm()` (loaded by default)

Generalized Linear Models: Stats package, `glm()`

Traditional Econometric Models: `plm` package

Mixed Modeling: `nlme` and `lme4` packages

Commonly Used Packages: Machine Learning

Most probably all you need is caret package.

Caret package is a wrapper for a host of classification and regression model training functions. It eases visualizations, data manipulation, and analytics among others. **It currently supports over 150 types of models.**

If you insist on using individual packages:


Classifiers: class package

Support Vector Machines: kernlab, e1071 packages

Clustering: Base package (kmeans(), hclust()), mclust package

Neural Networks: neuralnet package.

Questions



Outline

- 1 Introduction
- 2 R Console
- 3 Importing Data
- 4 Packages
- 5 Sample Analysis and Visualizations**
 - Descriptive Visualizations
 - Modeling
- 6 Reporting
- 7 Where to Go Next?

Subsection 1

Descriptive Visualizations

Motor Trends Dataset

We will use 1974 Motor Trend dataset. It has 32 observations and 11 variables.

- mpg: Miles per gallon
- cyl: Number of cylinders
- disp: Displacement
- hp: Horse Power
- drat: Rear axle ratio
- wt: Weight
- qsec: quarter mile time
- vs: V - S
- am: 0 automatic, 1 manual
- gear: Gears
- carb: Number of carburetors

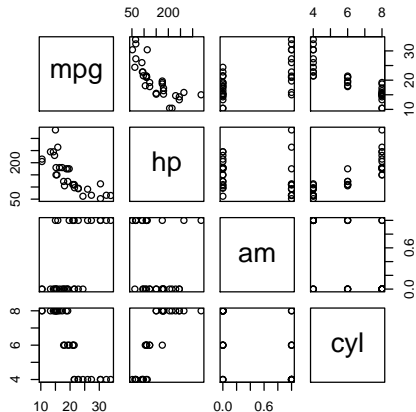
Motor Trends Dataset I

```
summary(mtcars)
```

```
##           mpg           cyl           disp           hp
##  Min.   :10.40   Min.   :4.000   Min.   : 71.1   Min.   : 52.0
##  1st Qu.:15.43   1st Qu.:4.000   1st Qu.:120.8   1st Qu.: 96.5
##  Median :19.20   Median :6.000   Median :196.3   Median :123.0
##  Mean   :20.09   Mean   :6.188   Mean   :230.7   Mean   :146.7
##  3rd Qu.:22.80   3rd Qu.:8.000   3rd Qu.:326.0   3rd Qu.:180.0
##  Max.   :33.90   Max.   :8.000   Max.   :472.0   Max.   :335.0
##
##      drat      wt      qsec      vs
##  Min.   :2.760   Min.   :1.513   Min.   :14.50   Min.   :0.0000
##  1st Qu.:3.080   1st Qu.:2.581   1st Qu.:16.89   1st Qu.:0.0000
##  Median :3.695   Median :3.325   Median :17.71   Median :0.0000
##  Mean   :3.597   Mean   :3.217   Mean   :17.85   Mean   :0.4375
##  3rd Qu.:3.920   3rd Qu.:3.610   3rd Qu.:18.90   3rd Qu.:1.0000
##  Max.   :4.930   Max.   :5.424   Max.   :22.90   Max.   :1.0000
##
##      am      gear      carb
##  Min.   :0.0000   Min.   :3.000   Min.   :1.000
##  1st Qu.:0.0000   1st Qu.:3.000   1st Qu.:2.000
##  Median :0.0000   Median :4.000   Median :2.000
##  Mean   :0.4062   Mean   :3.688   Mean   :2.812
##  3rd Qu.:1.0000   3rd Qu.:4.000   3rd Qu.:4.000
##  Max.   :1.0000   Max.   :5.000   Max.   :8.000
```



Motor Trends Dataset II

```
pairs(mtcars[, c("mpg", "hp", "am", "cyl")]) # Visualize Correlations
```



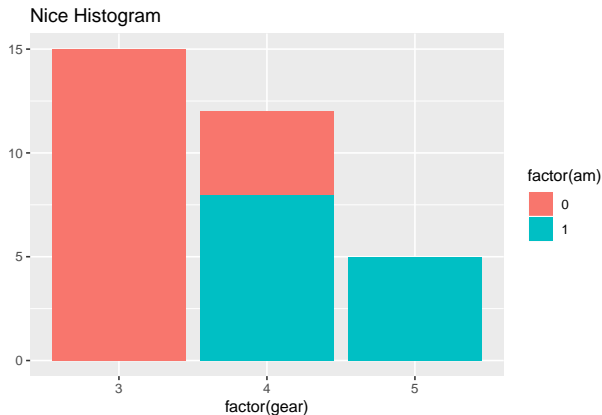
Histograms I

```
data(mtcars) # Load the Dataset
library(ggplot2) # Load the ggplot package
# Nr of cars by number of gears
ggplot(factor(gear), data = mtcars, geom = "bar")
```



Histograms II


```
# If we are interested in a third categorical variable vs:  
qplot(factor(gear), data=mtcars, geom="bar", fill=factor(am)) +  
ggtitle('Nice Histogram') # This is how you add a title
```



Scatter Plots I


```
# Two continuous variables
```

```
qplot(hp, mpg, data = mtcars, color = factor(cyl))
```




Scatter Plots II

```
# Add two more variables represented by color and size of points  
qplot(hp, mpg, data = mtcars, color = factor(cyl), size = disp, alpha = 0.09)
```



Bar Charts I

```
ggplot(mtcars, aes(x = factor(gear), y = mpg, fill = factor(vs)), color = factor(vs))  
  stat_summary(fun.y = mean, position = position_dodge(), geom = "bar")
```



Subsection 2

Modeling

Modeling

Developing a mathematical explanation for real life phenomena.

Useful in predicting outcomes (Y), and/or explaining the role of drivers (X).

Y = Outcome Variable, Dependent Variable...

X = Predictor, Independent variable...

Formula Interface

Pay close attention to how we specify the model.

R Model

$$Y \sim x_1 + x_2$$

This basic structure will remain constant across many R packages.

Nifty Tricks with Formula Interface

If you have lots of variables use the shortcut `'.'`.

$$\text{mpg} \sim .$$

You can do transformations on the fly, no need to create variables.

$$\log(\text{mpg} + 1) \sim .$$

Dummy variables.

$$\log(\text{mpg} + 1) \sim \text{factor}(\text{gears})$$

Regression

Y is continuous.

Prediction.


Explanation.

Linearity, Homoscedasticity, Normality, No Multicollinearity

Regression I

Regression visualized.

```
ggplot(mtcars, aes(x = hp, y = mpg)) + geom_point() + # Add a regression line  
geom_smooth(method = lm)
```



Multiple Regression Example

We will keep using motor trends data set.

Formula

$$Y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \epsilon$$

R Model

$$Y \sim x_1 + x_2$$

Regression I

```
# Let us estimate gas milage
reg_0 <- lm(mpg ~ hp + cyl + am, mtcars)
summary(reg_0)

##
## Call:
## lm(formula = mpg ~ hp + cyl + am, data = mtcars)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.864  -1.811  -0.158   1.492   6.013
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  30.88834    2.78422   11.094 9.27e-12 ***
## hp          -0.03688    0.01452   -2.540  0.01693 *
## cyl         -1.12721    0.63417   -1.777  0.08636 .
## am           3.90428    1.29659    3.011  0.00546 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.807 on 28 degrees of freedom
## Multiple R-squared:  0.8041, Adjusted R-squared:  0.7831
## F-statistic: 38.32 on 3 and 28 DF,  p-value: 4.791e-10
```

Regression II

```
## Diagnostics
library(car)
# Multicollinearity (vif>10 indicates MC)
vif(reg_0)

##          hp          cyl          am
## 3.900138 5.048209 1.647399

# Normality (p<0.05 indicates NN)
shapiro.test(reg_0$residuals)

##
## Shapiro-Wilk normality test
##
## data:  reg_0$residuals
## W = 0.98366, p-value = 0.8961

# Homoscedasticity (p<0.05 indicates Heteroscedasticity)
ncvTest(reg_0)

## Non-constant Variance Score Test
## Variance formula: ~ fitted.values
## Chisquare = 3.744471    Df = 1    p = 0.05298249
```

Regression III

```
# Access Fitted Values View first 3 predictions  
reg_0$fitted.values[1:3]
```

```
##      Mazda RX4 Mazda RX4 Wag      Datsun 710  
##      23.97302      23.97302      26.85433
```

```
# PREDICTING NEW DATA BASED ON MODEL  
newCar <- mtcars[3, ] # 3rd observation is Datsun 710  
newCar$am <- 0 # What if it was automatic?  
reg_0$fitted.values[3] # Previous estimate
```


```
## Datsun 710  
##      26.85433
```

```
predict(reg_0, newdata = newCar) # Datsun with automatic transmission
```

```
## Datsun 710  
##      22.95005
```


Regression IV

```
## Plot the residuals against observation  
qplot(data=mtcars, x = mpg, y = reg_0$residuals) + ##  
  stat_smooth(method = "lm", col = "red")
```



Regression V

```
## COMPARE MODELS
reg_1 <- lm(mpg ~ hp + cyl + am + wt, mtcars) # add weight
anova(reg_0, reg_1)

## Analysis of Variance Table
##
## Model 1: mpg ~ hp + cyl + am
## Model 2: mpg ~ hp + cyl + am + wt
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      28 220.55
## 2      27 170.00   1    50.555 8.0295 0.008603 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

# AIC of the model
AIC(reg_0)

## [1] 162.5849

AIC(reg_1)

## [1] 156.2536
```

Logistic Regression

Ordinary least squares regression was great for predicting continuous variables.

What if we are interested in predicting categories?

If there are only two categories (binary) logistic regression is a viable alternative.

Idea is similar to OLS regression. The only difference is we predict a constrained outcome (probability).

Logistic Regression

Dependent variable will be type (binary).

It is basically a regression with a binomial link function.

Formula

$$\log \left(\frac{p}{1-p} \right) = \beta_0 + \beta_1 x_1 + \epsilon$$

Logit I

```
logit_2 <- glm(am ~ mpg + drat + cyl, data = mtcars, family = "binomial")
summary(logit_2)
```


```
##
## Call:
## glm(formula = am ~ mpg + drat + cyl, family = "binomial", data = mtcars)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.58367  -0.31020  -0.03757   0.17972   1.75395
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -49.4548    24.1280  -2.050   0.0404 *
## mpg           0.6378     0.4266   1.495   0.1349
## drat          7.2595     3.2702   2.220   0.0264 *
## cyl           1.6115     1.0801   1.492   0.1357
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 43.23  on 31  degrees of freedom
## Residual deviance: 17.03  on 28  degrees of freedom
## AIC: 25.03
##
## Number of Fisher Scoring iterations: 7
```

Logit II

```
# VISUALIZE mpg - Transmission RELATION
ggplot(mtcars, aes(x = mpg, y = am)) +
  stat_smooth(method="glm", family="binomial", se=FALSE) #

## Warning: Ignoring unknown parameters: family
```

Logit III




Logit IV

```
## VISUALIZE Mpg - Transmission FOR DIFFERENT NUMBERS OF CYLINDERS

# Create a new dataset with varying number of cylinders and other variables
# fixed at mean levels.
mtcars2 <- data.frame(mpg = rep(10:30, 3), drat = mean(mtcars$drat), disp = mean(mtcars$disp),
  cyl = rep(c(4, 6, 8), 21))
# Predict probability of new data
mtcars2$prob <- predict(logit_2, newdata = mtcars2, type = "response")

# Plot the results
ggplot(mtcars2, aes(x = mpg, y = prob)) + geom_line(aes(colour = factor(cyl)),
  size = 1) # a different color for each category$
```


Logit V



Logit VI

```
## DIAGNOSTICS

# Let us compare predicted values to real values
mtcars$prob <- predict(logit_2, type = "response")
# Prevalence of Manual Transmission
mean(mtcars$am)

## [1] 0.40625


# Create predict variable
mtcars$pred <- 0
# If probability is greater than .6 (1-prevalence), set prediction to 1
mtcars[mtcars$prob > 0.6, "pred"] <- 1
```

Logit VII

```
## ROC CURVE

# Load the necessary library
library(pROC)
# Calculate the ROC curve using the predicted probability vs actual values
logit_2_roc <- roc(am ~ prob, mtcars)
# Plot ROC curve
plot(logit_2_roc)
```

Logit VIII




Logit IX

```
library(caret) # Needed for Confusion Matrix
confusionMatrix(table(mtcars[, c("am", "pred")]))
```

```
## Confusion Matrix and Statistics
##
##      pred
## am  0  1
##    0 18  1
##    1  3 10
##
##              Accuracy : 0.875
##              95% CI : (0.7101, 0.9649)
##    No Information Rate : 0.6562
##    P-Value [Acc > NIR] : 0.005004
##
##              Kappa : 0.7344
## Mcnemar's Test P-Value : 0.617075
##
##      Sensitivity : 0.8571
##      Specificity : 0.9091
##      Pos Pred Value : 0.9474
##      Neg Pred Value : 0.7692
##      Prevalence : 0.6562
##      Detection Rate : 0.5625
##      Detection Prevalence : 0.5938
##      Balanced Accuracy : 0.8831
##
##      'Positive' Class : 0
##
```

Questions



Outline

- 1 Introduction
- 2 R Console
- 3 Importing Data
- 4 Packages
- 5 Sample Analysis and Visualizations
 - Descriptive Visualizations
 - Modeling
- 6 Reporting**
- 7 Where to Go Next?

Markdown and R Markdown

Think HTML simplified.

Add (R) code into it.

Stir well with knitr

R Markdown Cheat Sheet

H1

H2

H3

Italic

****Bold****

****** Bold Italic******

```
``` { r, options }
```

```
R CODE HERE
```

```
```
```

Professional Looking Documents

All the documentation of this workshop has been prepared in R Studio.

Look into the Rmd documents to see how this was done.

`summary()` function does not provide the best looking model summaries.
Try the `texreg` package (do as I say, not as I do).

Use themes for your plots.


Other Reporting Alternatives

Sweave: \LaTeX

Beamer: \LaTeX presentations

Shiny: Web Applications

Questions



Outline

- 1 Introduction
- 2 R Console
- 3 Importing Data
- 4 Packages
- 5 Sample Analysis and Visualizations
 - Descriptive Visualizations
 - Modeling
- 6 Reporting
- 7 Where to Go Next?

Quo Vadis?

4 hours is not enough!

RTFM²: Documentation is your friend, read the vignettes and manuals.

For tutorials: [Institute For Digital Research and Education, UCLA](#)

For questions: StackExchange³


- [Cross Validated](#) specializes on Statistics
- [Stack Overflow](#) for R programming

R Programming is quirky, learn about efficient R programming.

²Read the Friendly Manual

³Word of caution, [pay attention to how you ask your questions](#). It is more important than you realize!

Questions



License



How I Learned to Stop Worrying and Love the R Console by **Irfan E Kanat** is licensed under a **Creative Commons Attribution 4.0 International License**. Based on a work at <http://github.com/iekanat/rworkshop>.